


## Comp3331 lab6

## Exercise 1: Setting up NS2 simulation for measuring TCP throughput

(submit the completed file `exercise2.tcl` and `throughput.plot` separately and answer the questions on throughput behaviour in the report)

Consider the topology shown in the following figure where bandwidth and delay for each link is shown.



You have been provided with a stub tcl file [exercise2.tcl](#). Your task is to complete the stub file so that it runs with ns and produces two trace files `tcp1.tr` and `tcp2.tr`, and `nam.out`. Check the animation for the simulation using `nam.out` file. Next write a script named "throughput.plot" (referenced from within `exercise2.tcl` in procedure `finish()`) to plot the throughput received by host n5 for two flows terminating at n5. Uncomment the line `(#exec gnuplot throughput.plot &)` to execute gnuplot. You have been provided with the throughput plot [TCPThroughput.png](#) produced by gnuplot for comparing your final output.

">>" in the stub file indicates that one (or more) lines need to be added. Remove the ">>" and insert the required code.

Consider the following traffic pattern for your simulation.

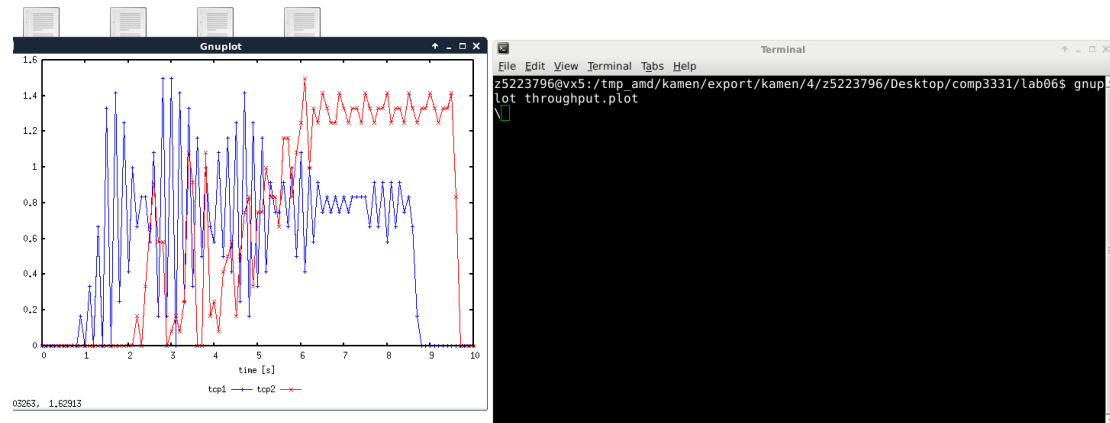
FTP/TCP Source n0 -> TCP Sink n5 : start time: 0.5 sec End time: 8.5 sec

FTP/TCP Source n3 -> TCP Sink n5 : start time: 2.0 sec End time: 9.5 sec

FTP/TCP Source n7 -> TCP Sink n0 : start time: 3.0 sec End time: 9.5 sec

FTP/TCP Source n7 -> TCP Sink n3 : start time: 4.0 sec End time: 7.0 sec

You have to submit your completed tcl file (exercise2.tcl) and the script (throughput.plot) for producing the throughput plot. Comment on the throughput behaviour observed in the simulation by answering the following questions.



**Question 1:** Why the throughput achieved by flow tcp2 is higher than tcp1 between time span 6 sec to 8 sec?

**Answer:**

Because RTT and number of flows competing of link bandwidth.

TCP1 will competing with TCP4 on link n1-n2, with TCP2 on link n2-n4

TCP2 has less RTT than tcp1, so TCP2 gets high share of bandwidth on link n2-n4. Thus the final throughput will higher.

**Question 2:** Why the throughput for flow tcp1 is fluctuating between time span 0.5 sec to 2 sec?

**Answer:**

At that time, TCP1 is doing the slow-start stage.

**Question 3:** Why is the maximum throughput achieved by any one flow capped at around 1.5Mbps?

**Answer:**

At the 0.5-2.0 sec, tcp1 is the only flow but it is in the slow-start stage and not achieve the maximum bandwidth . After TCP2 enter at 2.0 sec, it will compete with TCP1 and share the bandwidth with TCP1 which lead to TCP1 cannot get higher throughput.

## Exercise 2: Understanding IP Fragmentation

(Include in your report)

We will try to find out what happens when IP fragments a datagram by increasing the size of a datagram until fragmentation occurs. You are provided with a Wireshark trace file [IPfrag\\_trace](#) that contains trace of sending pings with specific payloads to 8.8.8.8. We have used ping with option ( -s option on Linux) to set the size of data to be carried in the ICMP echo request message. Note that the default packet size is 64 bytes in Linux (56 bytes data + 8 bytes ICMP header). Also note that Linux implementation for ping also uses 8 bytes of ICMP time stamp option leaving 48 bytes for the user data in the default mode. Once you have send a series of packets with the increasing data sizes, IP will start fragmenting packets that it cannot handle. We have used the following commands to generate this trace file.

Step 1: Ping with default packet size to the target destination as 8.8.8.8

```
ping -c 3 8.8.8.8
```

Step 2: Repeat by sending a set of ICMP requests with data of 2000.

```
ping -s 2000 -c 3 8.8.8.8
```

Step 3: Repeat again with data size set as 3500

```
ping -s 3500 -c 3 8.8.8.8
```

Load this trace file in Wireshark, filter on protocol field ICMP (you may need to clear the filter to see the fragments) and answer the following questions.

**Question 1:** Which data size has caused fragmentation and why? Which host/router has fragmented the original datagram? How many fragments have been created when data size is specified as 2000?

**Answer:**

No.	Time	Source	Destination	Protocol	Length	Info
5	0.000000000	192.168.1.103	8.8.8.8	ICMP	60	Echo (ping) request seq=0, ttl=64 (reply in 9)
6	0.000000000	8.8.8.8	192.168.1.103	ICMP	60	Echo (ping) reply seq=0, ttl=64 (request in 5)
9	0.000000000	192.168.1.103	8.8.8.8	ICMP	60	Echo (ping) request seq=0, ttl=64 (reply in 10)
10	0.000000000	8.8.8.8	192.168.1.103	ICMP	60	Echo (ping) reply seq=0, ttl=64 (request in 9)
13	0.000000000	192.168.1.103	8.8.8.8	ICMP	60	Echo (ping) request seq=0, ttl=64 (reply in 13)
14	0.000000000	8.8.8.8	192.168.1.103	ICMP	60	Echo (ping) reply seq=0, ttl=64 (request in 13)
17	0.000000000	192.168.1.103	8.8.8.8	ICMP	60	Echo (ping) request seq=0, ttl=64 (reply in 17)
18	0.000000000	8.8.8.8	192.168.1.103	ICMP	60	Echo (ping) reply seq=0, ttl=64 (request in 17)
20	0.000000000	192.168.1.103	8.8.8.8	ICMP	60	Echo (ping) request seq=0, ttl=64 (reply in 20)
21	0.000000000	8.8.8.8	192.168.1.103	ICMP	60	Echo (ping) reply seq=0, ttl=64 (request in 20)
41	0.000000000	192.168.1.103	8.8.8.8	ICMP	60	Echo (ping) request seq=0, ttl=64 (reply in 41)
42	0.000000000	8.8.8.8	192.168.1.103	ICMP	60	Echo (ping) reply seq=0, ttl=64 (request in 41)
44	0.000000000	192.168.1.103	8.8.8.8	ICMP	60	Echo (ping) request seq=0, ttl=64 (reply in 44)
45	0.000000000	8.8.8.8	192.168.1.103	ICMP	60	Echo (ping) reply seq=0, ttl=64 (request in 44)
50	0.000000000	192.168.1.103	8.8.8.8	ICMP	60	Echo (ping) request seq=0, ttl=64 (reply in 50)
51	0.000000000	8.8.8.8	192.168.1.103	ICMP	60	Echo (ping) reply seq=0, ttl=64 (request in 50)
57	0.000000000	192.168.1.103	8.8.8.8	ICMP	60	Echo (ping) request seq=0, ttl=64 (reply in 57)
58	0.000000000	8.8.8.8	192.168.1.103	ICMP	60	Echo (ping) reply seq=0, ttl=64 (request in 57)

2000 and 3500, Because the MTU is 1500 bytes.

192.168.1.103 fragmented the original datagram.

2 fragments have been created.

16	10.558043000	192.168.1.103	8.8.8.8	IPv4	1514	Fragmented IP protocol (proto=ICMP 1, off=0, ID=613d) [Reassembled in #1]
17	10.558045000	192.168.1.103	8.8.8.8	ICMP	562	Echo (ping) request id=0xd905, seq=0/0, ttl=64 (reply in 19)

**Question 2:** Did the reply from the destination 8.8.8.8. for 3500-byte data size also get fragmented? Why and why not?

**Answer:**

Yes, Since the MTU for the last link is 1500bytes and it has to put the last link to sender.

**Question 3:** Give the ID, length, flag and offset values for all the fragments of the first packet sent by 192.168.1.103 with data size of 3500 bytes?

20	192.168.1.103	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #41)
40	19.30871000	192.168.1.103	8.8.8.8	ICMP	1502	Echo (ping) request ID=0x000, seq=0, ttl=64 (reply in #4)
41	19.30871000	192.168.1.103	8.8.8.8	ICMP	1502	Echo (ping) request ID=0x000, seq=0, ttl=64 (reply in #4)
42	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #44)
43	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #44)
44	19.40802000	8.8.8.8	192.168.1.103	ICMP	1486	Echo (ping) reply ID=0x000, seq=0, ttl=64 (request in #4)
45	19.39621000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #47)
46	19.39621000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #47)
47	19.39621000	192.168.1.103	8.8.8.8	ICMP	1482	Echo (ping) request ID=0x000, seq=256, ttl=64 (reply in #8)
48	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #50)
49	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #50)
50	19.40802000	8.8.8.8	192.168.1.103	ICMP	1486	Echo (ping) reply ID=0x000, seq=256, ttl=64 (request in #7)
51	21.196617000	192.168.1.1	255.255.255.255	UDP	215	Source port: 36861 Destination port: 7437
52	21.403497000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #54)
53	21.403497000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #54)
54	21.403497000	192.168.1.103	8.8.8.8	ICMP	1482	Echo (ping) request ID=0x000, seq=2512, ttl=64 (reply in #7)
55	21.403497000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #57)
56	21.403497000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #57)

Internet Protocol Version 4, Src: 192.168.1.103, Dst: 192.168.1.103, Seq: 8.8.8.8 (8.8.8.8)

Version: 4

Header Length: 20 bytes

Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))

Total Length: 2000

Identification: 0x7a7b (31355)

Flags: 0x00 (None Fragments)

0 ..... = Reserved bit: Not set

0 ..... = Don't Fragment: Not set

1 ..... = More Fragments: Set

Fragment Offset: 0

Time to Live: 64

Id: 0x7a7b length: 1500 flag: 0x01 offset: 0

41	19.39671000	192.168.1.103	8.8.8.8	ICMP	1502	Echo (ping) request ID=0x000, seq=0, ttl=64 (reply in #4)
42	19.39671000	192.168.1.103	8.8.8.8	ICMP	1502	Echo (ping) request ID=0x000, seq=0, ttl=64 (reply in #4)
43	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #44)
44	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #44)
45	19.39621000	192.168.1.103	8.8.8.8	ICMP	1486	Echo (ping) reply ID=0x000, seq=0, ttl=64 (request in #4)
46	19.39621000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #47)
47	19.39621000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #47)
48	19.40802000	8.8.8.8	192.168.1.103	ICMP	1482	Echo (ping) request ID=0x000, seq=256, ttl=64 (reply in #8)
49	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #50)
50	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #50)
51	21.196617000	192.168.1.1	255.255.255.255	UDP	215	Source port: 36861 Destination port: 7437
52	21.403497000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #54)
53	21.403497000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #54)
54	21.403497000	192.168.1.103	8.8.8.8	ICMP	1482	Echo (ping) request ID=0x000, seq=2512, ttl=64 (reply in #7)
55	21.403497000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #57)
56	21.403497000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #57)

Internet 12, Src: 192.168.1.103, Dst: 192.168.1.103, Seq: 8.8.8.8 (8.8.8.8)

Internet Protocol Version 4, Src: 192.168.1.103, Dst: 192.168.1.103, Seq: 8.8.8.8 (8.8.8.8)

Version: 4

Header Length: 20 bytes

Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))

Total Length: 2000

Identification: 0x7a7b (31355)

Flags: 0x00 (None Fragments)

0 ..... = Reserved bit: Not set

0 ..... = Don't Fragment: Not set

1 ..... = More Fragments: Set

Fragment Offset: 0

Id: 0x7a7b length: 1500 flag: 0x01 offset: 1480

42	19.39571000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #41)
43	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #44)
44	19.40802000	8.8.8.8	192.168.1.103	ICMP	1486	Echo (ping) reply ID=0x000, seq=0, ttl=64 (request in #4)
45	19.39621000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #47)
46	19.39621000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #47)
47	19.39621000	192.168.1.103	8.8.8.8	ICMP	1482	Echo (ping) request ID=0x000, seq=256, ttl=64 (reply in #8)
48	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #50)
49	19.40802000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #50)
50	19.40802000	8.8.8.8	192.168.1.103	ICMP	1486	Echo (ping) reply ID=0x000, seq=256, ttl=64 (request in #7)
51	21.196617000	192.168.1.1	255.255.255.255	UDP	215	Source port: 36861 Destination port: 7437
52	21.403497000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #54)
53	21.403497000	192.168.1.103	8.8.8.8	IPV4	1514	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #54)
54	21.403497000	192.168.1.103	8.8.8.8	ICMP	1482	Echo (ping) request ID=0x000, seq=2512, ttl=64 (reply in #7)
55	21.403497000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #57)
56	21.403497000	8.8.8.8	192.168.1.103	IPV4	1482	Fragmented IP protocol (proto=TCP 1, offset=0, ID=0x7a7b) (Reassembled in #57)

Internet 12, Src: 192.168.1.103, Dst: 192.168.1.103, Seq: 8.8.8.8 (8.8.8.8)

Internet Protocol Version 4, Src: 192.168.1.103, Dst: 192.168.1.103, Seq: 8.8.8.8 (8.8.8.8)

Version: 4

Header Length: 20 bytes

Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))

Total Length: 568

Identification: 0x7a7b (31355)

Flags: 0x00 (None Fragments)

0 ..... = Reserved bit: Not set

0 ..... = Don't Fragment: Not set

1 ..... = More Fragments: Not set

Fragment Offset: 2960

Id: 0x7a7b length: 568 flag: 0x00 offset: 2960

**Question 4:** Has fragmentation of fragments occurred when data of size 3500 bytes has been used? Why and why not?

No fragmentation from 8.8.8.8 to 192.168.1.103, we have three fragments, no fragmentation of fragments. And we can not sure about the fragmentation from 192.168.1.103 to 8.8.8.8, because the reassembled only occur the destination.

**Question 5:** What will happen if for our example one fragment of the original datagram from 192.168.1.103 is lost?

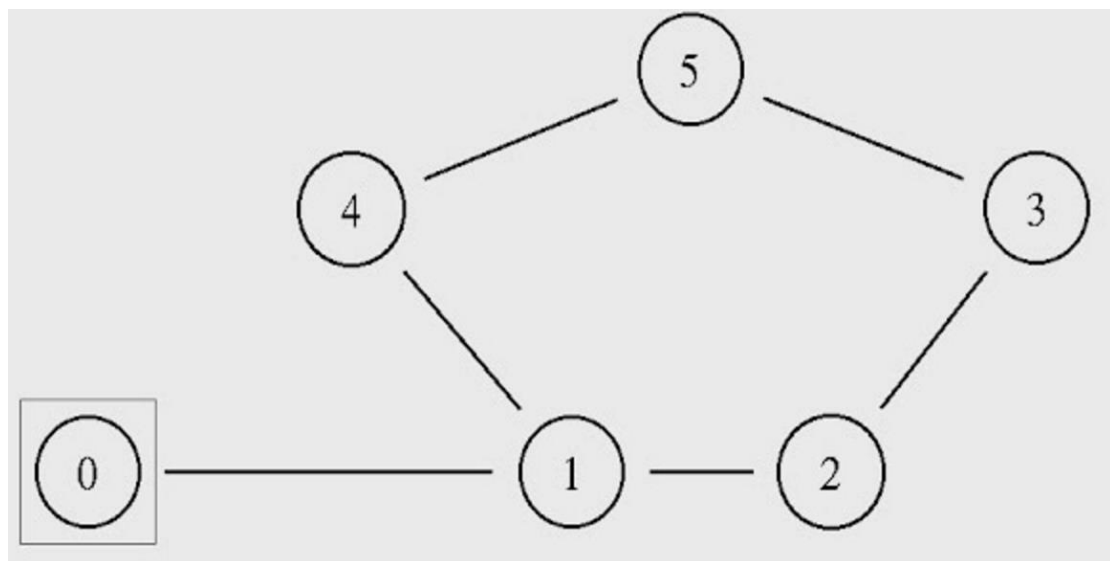
Fragment will be incomplete and receiver will discard.

## Exercise 3: Understanding the Impact of Network Dynamics on Routing

(include in your report)

In this exercise, we will observe how routing protocols react when network conditions change (e.g., a network link fails) using a ns-2 simulation.

The provided script, [tp\\_routing.tcl](#) takes no arguments and generates the network topology shown in the figure below.



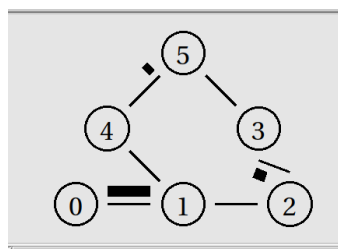
You can run the simulation with the following command:

```
$ns tp_routing.tcl
```

Step 1: Run the script and observe the NAM window output.

**Question 1:** Which nodes communicate with which other nodes? Which route do the packets follow? Does it change over time?

**Answer:**



node0->node5 (0->1->4->5)

node2->node5 (2->3->5) Does not change.

**Note:** You can also answer the above question by examining the simulation setting in the script file.

Step 2: Modify the script by uncommenting the following two lines (line No 84 and 85):

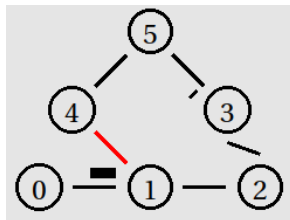
```
$ns rtmodel-at 1.0 down $n1 $n4  
$ns rtmodel-at 1.2 up $n1 $n4
```

Step 3: Rerun the simulation and observe the NAM window output.

**NOTE:** Ignore the NAM syntax warnings on the terminal. These will not affect the simulation.

**Question 2:** What happens at time 1.0 and at time 1.2? Does the route between the communicating nodes change as a result of that?

**Answer:**



At time 1.0, link n1 -> n4 goes down. Node 0 can not reach node 5, and packets wait at node 1 but node 2 still can go node 5.

At time 1.2, link n1->n4 goes up. Node 0 can reach node 5. Node 1 can go node 4.

No influence on node 2 and node 5.

Step 4: The nodes in the simulation above use a static routing protocol (i.e., preferred routes do not change over time). We are going to change that, so that they use a Distance-Vector routing protocol. Modify the script and uncomment the following line (Line No 16) before the definition of the `finish` procedure.

```
$ns rtproto DV
```

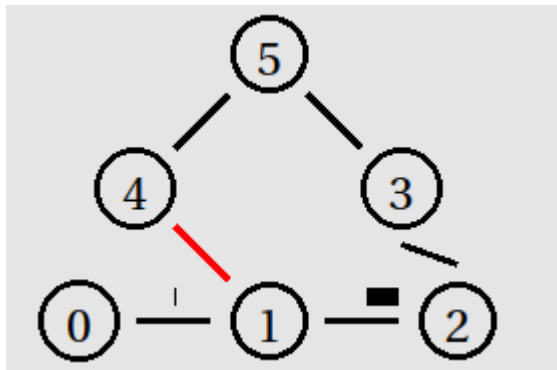
Step 5: Rerun the simulation and observe the NAM window output.

**Question 3:** Did you observe any additional traffic as compared to Step 3 above? How does the network react to the changes that take place at time 1.0 and time 1.2 now?

**Answer:**

Yes, When the link n1 -> n4 goes down, Node 0 to node 5 will use another path

0->1->2->3->5 and when n1-n4 goes up, it will 0-1-4-5



Step 6: Comment the two lines (Lines 84 and 85) that you had added to the script in Step 2 and uncomment the following line ( Line 87) instead:

```
$ns cost $n1 $n4 3
```

Step 7: Rerun the simulation and observe the NAM window output.

**Question 4:** How does this change affect the routing? Explain why.

**Answer:**

It increase the n1-n4 cost to 3 and the total cost for route 0-1-4-5 is 5

And 0-1-2-3-5 will be 4. We choose lower cost route. So for node 0 to node 5 will use route 0-1-2-3-5

Step 8: Comment line 87 and Uncomment the following lines (Lines 89 and 90):

```
$ns cost $n1 $n4 2
```

```
$ns cost $n3 $n5 3
```

and uncomment the following (Line 29), which is located right after the finish procedure definition:

```
Node set multiPath_1
```

Step 9: Rerun the simulation and observe the NAM window output.

**Question 5:** Describe what happens and deduce the effect of the line you just uncommented.

**Answer:**

For node 0 to node 5, route 0-1-4-5 is still the lower cost (4) route. But for node 2 to node 5, route 2-3-5 and 2-1-4-5 has the same cost 4 and it will split equally on these two routes.

